

Hypermedia In Action

1. Reintroducing Hypermedia

This chapter covers

- A reintroduction to the core concepts of hypermedia
- Why you might choose hypermedia over other approaches
- How hypermedia can be used to build modern web applications

Hypermedia is a universal technology today, nearly as common as electricity. Billions of people use a hypermedia-based systems every day, mainly by interacting with the *HyperText Markup Language (HTML)* being exchanged via the *HyperText Transfer Protocol (HTTP)* by using a Web Browser connected to the World Wide Web. People use these systems to get their news, check in on friends, buy things online, play games, send emails and so forth: the variety and sheer number of online services is truly astonishing.

And yet, despite this ubiquity, hypermedia itself is a strangely under-explored concept, left mainly to specialists. Yes, you can find a lot of tutorials on how to author HTML, create links and forms, etc. But it is rare to see a discussion of HTML as a *hypermedia*. This is in sharp contrast with the early web development era, when concepts like *Representational State Transfer (REST)* and *Hypermedia As The Engine of Application State (HATEOAS)* were constantly discussed and debated among developers.

It is sad to say, but in some circles today HTML is viewed almost resentfully: it is considered an awkward, legacy user interface description language that must be used build Javascript-based applications, simply because it happens to be there, in the browser.

This as a shame, and we hope that with this book we can convince you that the hypermedia architecture is not simply a piece of legacy technology that we have to begrudgingly deal with. Instead, we aim to show that it is a tremendously innovative, flexible and *simple* way to build robust distributed systems. Not only that, but this hypermedia approach deserves a seat at the table when you, a web developer, are considering the architecture of your next online software system.

1.1. So, What Is Hypermedia?

The English prefix "hyper-" comes from the Greek prefix " $\overline{\alpha}\pi\epsilon\rho-$ " and means "over" or "beyond"... It signifies the overcoming of the previous linear constraints of written text.

— Wikipedia, <https://en.wikipedia.org/wiki/Hypertext>

Right. So what is hypermedia? Simply, it is a media, for example a text, that includes non-linear branching from one location to another, via, for example, hyperlinks embedded directly in the media.

You are probably more familiar with the term *hypertext*, from whose Wikipedia page the above quote is taken. Hypertext is a sub-set of hypermedia and much of this book is going to discuss how to build modern web applications with HTML, the HyperText Markup Language.

However, even when working with applications built mainly in HTML, there are nearly always other medias involved: images, videos and so forth. This makes *hypermedia* a more appropriate term for discussing applications built in this manner. We will use the term hypermedia for most of this book in order to capture this more general concept.

1.1.1. HTML

In the beginning was the hyperlink, and the hyperlink was with the web, and the hyperlink was the web. And it was good.

— Rescuing REST From the API Winter, <https://intercoolerjs.org/2016/01/18/rescuing-rest.html>

Before we discuss the more general concepts of hypermedia, let's take a brief look at a concrete, familiar example of it: HTML.

HTML is the most widely used hypermedia in existence, and this book naturally assumes that the reader has a reasonable familiarity with it. You don't need to be an HTML or CSS expert to understand the code in this book, but the better you understand the core tags and concepts of both HTML, the more you will get out of this book.

Now, let's consider the two defining elements of hypermedia in HTML: the anchor tag (which produces a hyperlink) and the form tag.

Here is a simple anchor tag:

Listing 1. 1. A Simple Hyperlink

```
<a href="https://www.manning.com/">  
    Manning Books  
</a>
```

In a typical browser, this tag would be interpreted to mean: "Show the text 'Manning Books' in manner indicating that it is clickable and, when the user clicks on that text, issue an HTTP GET to the url <https://www.manning.com/>. Take the resulting HTML content in the body of the response and use it to replace the entire screen in the browser as a new document."

This is the main mechanism we use to navigate around the web today, and it is a canonical example of a hypermedia link, or a hyperlink.

So far, so good. Now let's consider a simple form tag:

Listing 1. 2. A Simple Form

```
<form action="/signup" method="post">
  <input type="text" name="email" placeholder="Enter Email To Sign Up..."/>
  <button>Sign Up</button>
</form>
```

This bit of HTML would be interpreted by the browser roughly as: "Show a text input and button to the user. When the user submits the form by clicking the button or hitting enter in the input, issue an HTTP POST request to the path '/signup' on the site that served the current page. Take the resulting HTML content in the response body and use it to replace the entire screen in the browser."

I am omitting a few details and complications here: you also have the option of issuing an HTTP **GET** with forms, the result may *redirect* you to another URL and so on, but this is the crux of the form tag.

Here is a visual representation of these two hypermedia interactions:

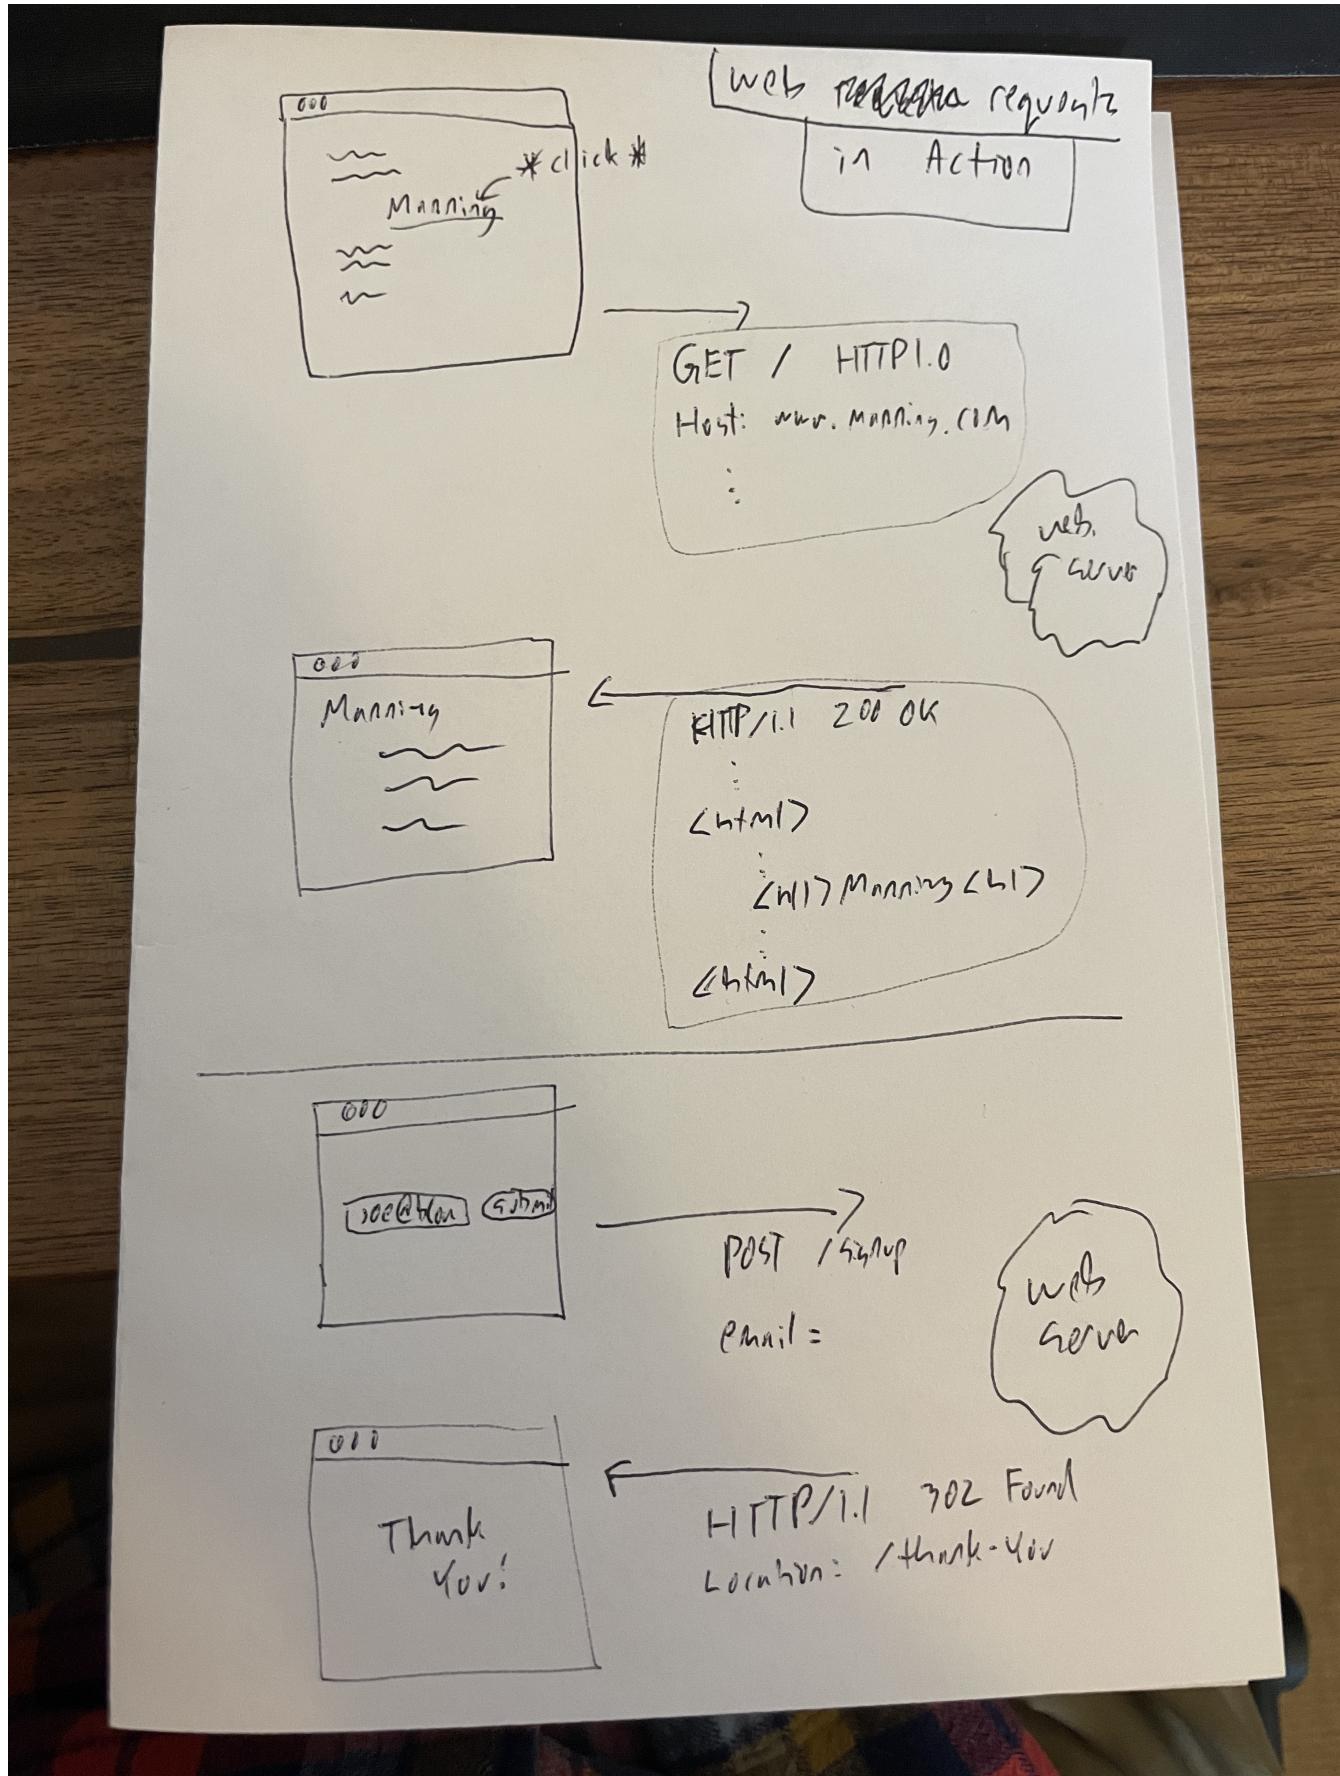


Figure 1. 1. HTTP Requests In Action

As someone interested in web development, the above diagram should look very familiar to you, perhaps even boring. But, despite its familiarity, consider the fact that the two above mechanisms are the *only* easy ways to interact with a server via HTML. That's barely anything at all! And yet,

armed with only these two tools, the early web was able to grow exponentially and offer a staggeringly large amount of functionality to an even more staggeringly large number of people!

This is strong evidence of the power of hypermedia. Even today, in a web development world increasingly dominated by large JavaScript-centric front end frameworks, many people choose to simply use vanilla HTML to achieve their goals and are perfectly happy with the results.

With just these two little tags, hypermedia manages to pack a heck of a punch!

1.1.2. So What Isn't Hypermedia?

So we've looked at the two ways to interact with a server via HTML directly. Now let's consider another approach to interacting with a server:

Listing 1.3. Javascript

```
<button onclick="fetch('/api/v1/contacts').then(response => response.json()).then(data => updateTable(data))">  
    Fetch Contacts  
</button>
```

Here we have a button element in HTML that executes some JavaScript when it is clicked. That JavaScript will issue a `GET` request to `/api/v1/contacts` using the `fetch()` API, a popular API for issuing an "Asynchronous JavaScript and XML", or AJAX request. An AJAX request is like a normal HTTP request in many ways, but it is issued "behind the scenes" by the browser: the user does not see a request indicator and it is up to the JavaScript code to deal with the response.

Despite the name, the response to this request will almost certainly be in the JavaScript Object Notation (JSON) format rather than XML. (That is a long story!)

The response to this request might look something like this:

Listing 1.4. JSON

```
{  
  "id": 42,  
  "email" : "json-example@example.org"  
}
```

The JavaScript code above converts this JSON text received from the server into a Javascript object (which is very easy when using the JSON notation) which is then handed off to the `updateTable()` method. The `updateTable()` method, not shown here to keep things simple, would then update the UI based on the data that has been received from the server.

What we want to stress at this point of the book is that this server interaction is *not* using hypermedia. The JSON API being used here does not return a hypermedia response. This is, rather, a *Data API*, returning simple, plain old domain data in JSON format. It is up to the code in the `updateTable()` method to understand how to turn this plain old data into HTML, which would typically be done via some sort of client-side templating library.

This bit of javascript is the beginnings of what has come to be called a Single Page Application (SPA): the application is no longer navigating between pages using hypermedia. Instead, we are, within a single page, exchanging *data* with the server and updating the content within that page.

Of course, today, the vast majority of web applications adopt far more sophisticated frameworks for managing the user interface than this simple example, libraries like React, Angular, Vue.js, etc. With these more complex frameworks you typically work with a much more elaborate client-side model (that is, JavaScript objects stored locally in the browser's memory that represent the model of your application.) You update these JavaScript objects in memory and then you allow the UI to "react" to those changes via infrastructure baked into the framework itself. (This is where the term "Reactive" programming comes from.)

In this approach, you, the developer, do not interact with hypermedia much at all. You use it to build your user interface, but the anchor tag is de-emphasized and forms become mere data collection mechanisms. Neither interact with the server in their native language of HTML, and instead become user interface elements that drive local interactions with the in memory domain model.

So modern SPAs are much more complex than the basic JavaScript example above. However, at the level of a *network architecture*, these more sophisticated frameworks are essentially equivalent to our simple example: they exchange plain data via JSON with the server rather than exchanging hypermedia.

1.2. Why Use Hypermedia?

The emerging norm for web development is to build a React single-page application, with server rendering. The two key elements of this architecture are something like:

1. The main UI is built & updated in JavaScript using React or something similar.
2. The backend is an API that that application makes requests against.

This idea has really swept the internet. It started with a few major popular websites and has crept into corners like marketing sites and blogs.

— Tom MacWright, <https://macwright.com/2020/05/10/spa-fatigue.html>

Tom is correct: JavaScript-based Single Page Applications have taken the web development world by storm, offering a far more interactive and immersive experience than the old, gronky, web 1.0 HTML-based application could. Some SPAs are even able to rival native applications in their user experience and sophistication.

So, why on earth would you abandon this new, increasingly standard (just do a job search for reactjs!) approach for an older and less discussed one like hypermedia?

Well, it turns out that, even in its original form, the hypermedia architecture has a number of

advantages when compared with the JSON/Data API approach:

- It is an extremely simple approach to building web applications
- It survives network outages and changes relatively well
- It is extremely tolerant of content and API changes (in fact, it thrives on them!)

As someone interested in web development, these advantages no doubt sound appealing to you. The first and last one, in particular, address two pain points in modern web development:

- Front end infrastructure has become extremely complex (sophisticated might be the nice way of saying it!)
- API churn is a huge pain for many applications

Taken together, these two problems have become known as "Javascript Fatigue": a general sense of exhaustion with all the hoops that are necessary to jump through to get anything done on the web.

And it's true: the hypermedia architecture *can* help cure Javascript Fatigue. But you may reasonably be wondering: so, if hypermedia is so great and can address these problems so obviously in the web development industry, why has it been largely abandoned by web developers today? After all, web developers are a pretty smart lot. Why wouldn't they use this obvious, native web technology?

In our opinion there are two related reasons for this somewhat strange state of affairs. The first is this: hypermedia (and HTML in particular) hasn't advanced much *since the late 1990s* as hypermedia. Sure, lots of new features have been added to HTML, but there haven't been *any* new ways to interact with a server via pure HTML added in over two decades! HTML developers are still working with only anchor tags and forms, and can only issue **GET** and **POST** requests.

This somewhat baffling lack of progress leads immediately to the second and more practical reason that hypermedia has been abandoned: as the interactivity and expressiveness of HTML remained frozen in time, the technology world marched on, demanding more and more interactive web applications. JavaScript, coupled to data-oriented JSON APIs, stepped in as a way to provide these interactive features in web applications to end users. It was this, the *user experience* that really drove the web developer community over to the JavaScript-heavy Single Page Application approach.

This is unfortunate, and it didn't have to be this way. There is nothing *intrinsic* to the idea of hypermedia that prevents a richer, more expressive interactivity model. Rather than abandoning the hypermedia architecture, the industry could have kept pushing it forward and enabling more and more interactivity *within* that original, hypermedia model of the web. If history had worked out that way, perhaps we could have retained much of the simplicity of the original web while still providing better user experiences.

1.2.1. A Hypermedia Comeback?

So, for many developers today working in an industry dominated by JavaScript and SPA frameworks, hypermedia has become an afterthought, if it is thought of at all. You simply can't get the sort of modern interactivity out of HTML, the hypermedia we all use day to day, necessary for today's modern web applications.

But, what if history *had* worked out differently?

What if HTML, instead of stalling as a hypermedia, had continued to develop, adding new mechanisms for exchanging hypermedia with servers and increasing its general expressiveness?

What if this made it possible to build modern web applications within the original, hypermedia-oriented and REST-ful model that made the early web so powerful, so flexible, so... fun? Could hypermedia be a legitimate architecture to consider when developing a new web application?

The answer is yes and, in fact, in the last decade, some alternative front end libraries (ironically, written in JavaScript!) have arisen that attempt to do exactly this. These libraries use JavaScript not as a *replacement* for the hypermedia architecture, but rather use it to augment HTML itself *as a hypermedia*.

These *hypermedia-oriented* libraries re-center the hypermedia approach as a viable and, indeed, excellent architectural choice for your next web application.

1.2.2. Hypermedia-Oriented Javascript Libraries

In the web development world today there is a debate going on between the SPAs approach and what are now being called "Multi-Page Applications" or MPAs. MPAs are, usually, just the old, traditional way of building web applications and thus are, by their nature, hypermedia oriented, if a bit clunky. Despite this clunkiness, some web developers have become so exasperated at the complexity of SPA applications they have decided to go back to this older way of building things and just accepting the less interactive nature inherent to plain HTML.

Some thought leaders in web development, such as Rich Harris, creator of svelte.js, propose a mix of the two styles. Harris calls this approach to building web applications "Transitional", in that it attempts to mix both the old MPA approach and the newer SPA approach in a coherent whole.

Again, the crux of the tradeoffs between SPAs and MPAs is the *user experience* or interactivity of the application, this is typically the driving decision when choosing one approach versus the other for an application or, in the case of Transitional Web Applications, for a feature.

However, by adopting a hypermedia oriented library, it turns out that this interactivity gap closes dramatically between the MPA and SPA approach. A hypermedia oriented library allows you to make the decision based on other considerations, such as overall system complexity.

One such hypermedia oriented library is htmx, created by the authors of this book. htmx will be the focus of much (but not all!) of the remainder of this book, and we hope to show you that you can, in fact, create many common "modern" UI features in a web application entirely within the hypermedia model. Not only that, but it is refreshingly fun and simple to do so!

When building a web application with htmx and other hypermedia oriented libraries the term Multi-Page Application applies *roughly*, but it doesn't really capture the crux of the application architecture. htmx, as you will see, does not need to replace entire pages and, in fact, an htmx-based application can reside entirely within a single page. (We don't recommend this practice, but it is certainly doable!)

We rather like to emphasize the *hypermedia* aspect of both the older MPA approach and the newer

htmx-based approach. Therefore, we use the term *Hypermedia Driven Applications (HDAs)* to describe both. This clarifies that the core distinction between these approaches and the SPA approach *isn't* the number of pages in the application, but rather the underlying *network* architecture.

What would the htmx and, let us say, the HDA equivalent of the JavaScript-based SPA-style button we discussed above look like?

It might look like this:

Listing 1.5. an htmx implementation

```
<button hx-get="/contacts" hx-target="#contact-table">  
    Fetch Contacts  
</button>
```

As with the JavaScript example, we see that this button has been annotated with some attributes. However, in this case we do not have any imperative scripting going on. Instead, we have *declarative* attributes, much like the `href` attribute on anchor tags and the `action` attribute on form tags. The `hx-get` attribute tells htmx: "When the user clicks this button, issue a `GET` request to `/contacts`"`. The `'hx-target'` attribute tells htmx: "When the response returns, take the resulting HTML and place it into the element with the id `'contact-table`"`.

I want to emphasize here that the response here is expected to be in *HTML format*, not in JSON. This means that htmx is exchanging hypermedia with the server, and thus the interaction is still firmly within this original hypermedia model of the web. htmx is adding browser functionality via JavaScript, but that functionality is *augmenting* HTML as a hypermedia, rather than *replacing* the network model with a Data-oriented JSON API.

So, despite perhaps looking superficially similar to one another, it turns out that this htmx example and the JavaScript-based example are extremely different architectures and approaches to web development. And, similarly, the htmx/HDA approach is extremely different from the SPA approach.

This may seem all well and good: a contrived little demo of a simple tool that maybe makes HTML a bit more expressive. But surely this is just a toy! It can't scale up to large, complex modern web applications, can it?

In fact, it can: just as the original web scaled up confoundingly well via hypermedia, due to the simplicity this approach it can often scale extremely well with your application needs. And, despite its simplicity, I think you will be surprised at just how much we can accomplish in creating modern, sophisticated user experiences in your web applications.

1.3. When should You Use Hypermedia?

Even if you decide not to use something like htmx and just accept the limitations of plain HTML, there are times when it, and the hypermedia architecture, is worth considering for your project:

Perhaps you are building a web application that doesn't *need* a huge amount of user-experience

innovation. These are very common and there is no shame in that! Perhaps your application adds its value on the server side, by coordinating users or by applying sophisticated data analysis. Perhaps your application adds value by simply fronting a well designed database with simple Create-Read-Update-Delete (CRUD) operations. Again, there is no shame in this!

In any of these later cases, using a hypermedia approach would likely be a great choice: the interactivity needs of these applications are not off the charts, and much of the value lives on the server side, rather than on the client side. They are all amenable to "large-grain hypermedia data transfers", exactly what the web was designed to do.

By adopting the hypermedia approach for these applications, you will save yourself a huge amount of client-side complexity: there is no need for client-side routing, for managing a client side model, for hand-wiring in javascript logic. You will be able to focus your efforts on your server, where your application is actually adding value.

And by layering htmx or another hypermedia-oriented library on top of this approach, you can address many of the usability issues of it with finer-grained hypermedia transfers. This opens up a whole slew of new user interface and experience possibilities, and will be a focus of later chapters.

1.4. When shouldn't You Use Hypermedia?

But, that being said, there are cases where hypermedia is not the right choice. What would a good example be?

One example that springs to mind is an online spreadsheet application, where updating one cell could have a large number of cascading changes that need to be made on every keystroke. In this case, we have a highly inter-dependent user interface without clear boundaries as to what might need to be updated given a particular change. Additionally, introducing a server round-trip on every cell change would bog performance down terribly. This is simply not a situation amenable to that "large-grain hypermedia data transfer" approach. We would heartily recommend using JavaScript-based infrastructure for building an application like this!

However, perhaps this online spreadsheet application also has a settings page. And perhaps that settings page is amenable to the hypermedia approach. If it is simply a set of relatively straightforward forms that need to be persisted to the server, the chances are high that hypermedia would, in fact, work great for this part of the app.

And, by adopting hypermedia for that part of your application, you could simplify this part of the application quite a bit. You can then save more of your *complexity budget* for the core, complicated spreadsheet logic, keeping the simple stuff simple.

What Is A Complexity Budget?

Any software project has a complexity budget, explicit or not: there is only so much complexity a given development team can tolerate and every new feature and implementation choice adds at least a bit more to the overall complexity of the system.

What is particularly nasty about complexity is that it appears to grow exponentially: one day you can keep the entire system in your head and understand the ramifications of a particular change, and a week later the whole system seems intractable. Even worse, efforts to help control complexity, such as introducing abstractions or infrastructure to manage the complexity, often end up making things even more complex. Truly, the job of the good software engineer is to keep complexity under control.

The surefire way to keep complexity down is also the hardest: say no. Pushing back on feature requests is an art and, if you can learn to do it well, making people feel like *they* said no, you will go far.

Sadly this is not always possible: some features will need to be built. At this point the question becomes: "what is the simplest thing that could possibly work?" Understanding the possibilities available in the hypermedia approach will give you another tool in that "simplest thing" tool chest.

This brings up two important points:

First, even the most hard core SPA application is, at some level, a "Transitional" web application: there is always a bootstrap page that gets the app started that is served via, wait for it, hypermedia! So you are already using the hypermedia approach when you build web applications, whether you think so or not.

Second, the hypermedia approach, in both its simple, "vanilla" HTML form and in its more sophisticated htmx form, can be adopted incrementally: you don't need to use this approach for your entire application. You can, instead, adopt it where it makes sense. Or, alternatively, you might flip this around and make hypermedia your default approach and only reach for the more complicated JavaScript-based solutions when necessary. We favor this approach in general as the ideal way to minimize your web applications complexity.

1.5. Summary

In this chapter, you have been reintroduced the concept of hypermedia and, I hope, you have a better understanding of what it is, and how it is an integral aspect of the web. We also talked about what is *not* hypermedia, focusing on a simple Data-Oriented JSON API. Here are the important points to remember:

- Hypermedia is a unique architecture for building web applications
- Using Data APIs, which is very common in today's web development world, is very dramatically different than the hypermedia approach
- Hypermedia lost out to SPAs & Data APIs due to interactivity limitations, not due to fundamental

limitations of = the concept

- There is an emerging class of front-end libraries, htmx being one, that recenter hypermedia as the core technology for web development and address these interactivity limitations
- These libraries make Hypermedia Driven Applications (HDAs) a more compelling choice for a much larger set of online applications

Before we get into the practical details of implementing a modern Hypermedia Driven Application, in the next chapter we will take a bit of time to make an in-depth study of one of the most influential documents ever written on the hypermedia architecture of the web, the famous Chapter 5 of Roy Fielding's PhD dissertation and the two concepts it defined and codified: REpresentational State Transfer (REST) & Hypermedia As The Engine of Application State (HATEOAS).